

# Milestone 3

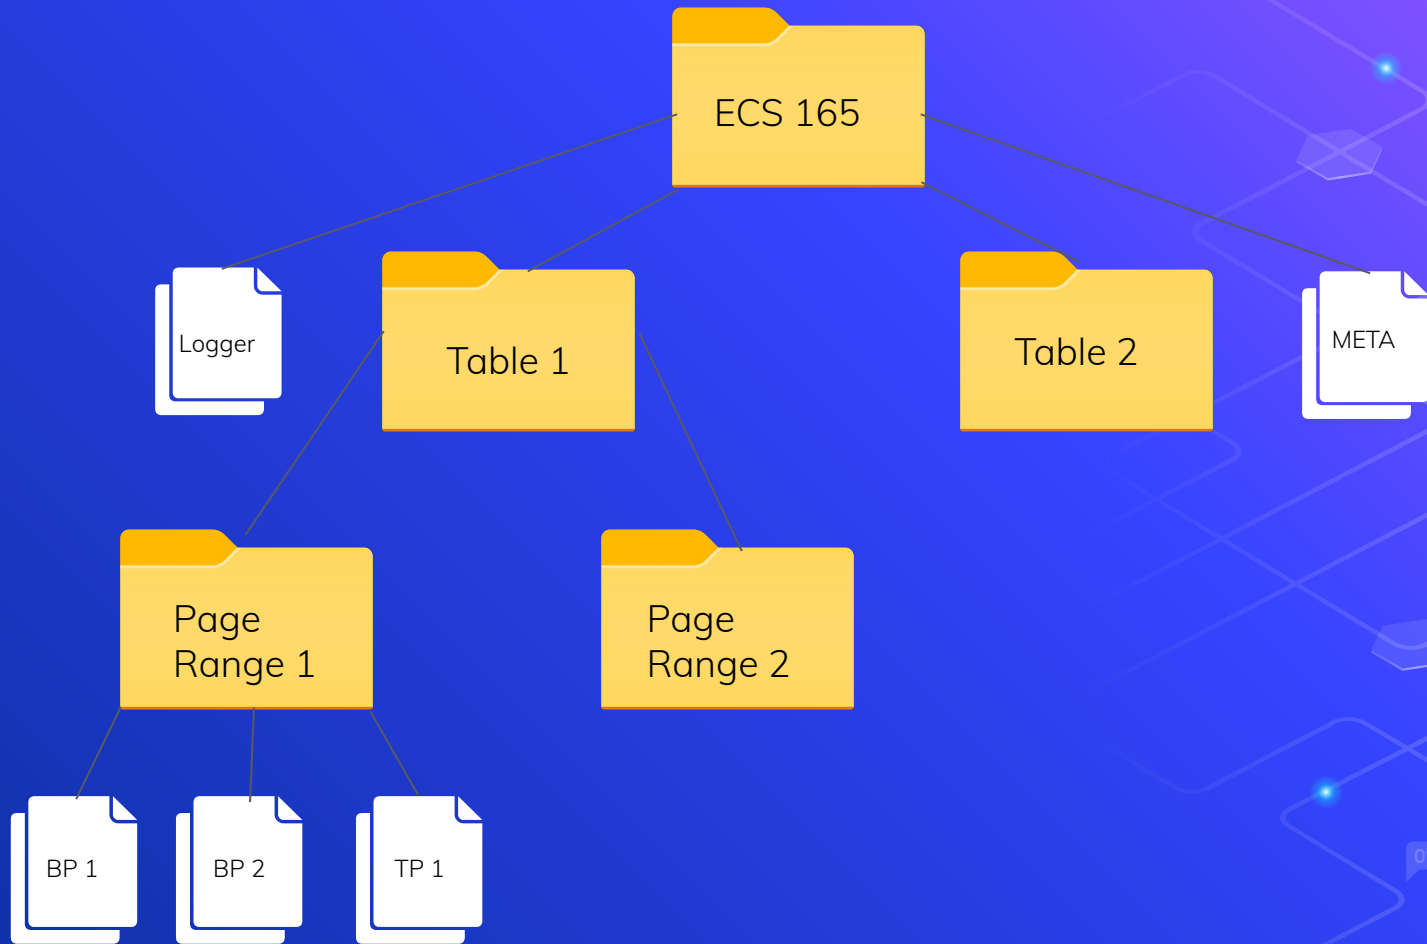
by The B Team



# Architecture

Minor Changes






# Transactional Semantics

- ⬡ Atomicity
- ⬡ Isolation



# Goal : Atomicity

- Single unit transactions
- All or nothing
  - Any fails => abort()
  - No fails => commit()
- Transaction Failure:
  - Query fails



# Transaction

- Multiple query calls
- run( ) , abort( ), and commit( )

Transaction 1
insert ( 91231, 23, 2 )
update ( 91231, 26, 3 )
insert ( 91233, 77 , 2 )




# Transaction Run()

- Execute in order
- Write to log
- Unsuccessful query -> halt execution
- Finishes without aborting



# Transaction Abort()

- Reads through the logs
- Abort Inserts deletes that record.
- Abort Updates removes the aborted updates and points to a new tail unaborted values.

Transaction 1		
insert ( 91231, 23, 2 )		
update ( 91231, 26, 3 )		
insert ( 91233, 77 , 2 )		



Abort()



# Transaction Commit()

- ⬡ Release all locks for a transaction number
- ⬡ Write transaction to disk

Transaction 1	
insert ( 91231, 23, 2 )	✓
update ( 91231, 26, 3)	✓
insert ( 91233, 77 , 2)	✓




# Logger

- ❖ Log each Query
- ❖ Transaction number, query type, key, and columns
- ❖ Undo Changes
- ❖ DB failure (power failure, etc)



# Concurrency Control

+ Multithreading



# Goal : Isolation

- Concurrent Transactions
- Lock manager
- Treat as sequential



# Transaction Worker

- ❖ List of transactions
- ❖ Run on own thread
- ❖ Each thread runs concurrently, wait for all threads to finish before we close



# Lock Manager

- ⬡ Read Lock
- ⬡ Write Lock
- ⬡ Abort if locked



# Latches

- ⬡ Latch critical points
- ⬡ Bufferpool locked during read/write operations
- ⬡ Insert locked



# Thank You!





# Extra resources






SlidesCarnival icons are editable shapes.

This means that you can:

- Resize them without losing quality.
- Change fill color and opacity.
- Change line color, width and style.


Isn't that nice? :)

Examples:



Find more icons at  
[slidescarnival.com/extra-free-resources-icons-and-maps](https://slidescarnival.com/extra-free-resources-icons-and-maps)

## Diagrams and infographics



You can also use any emoji as an icon!  
And of course it resizes without losing quality.

How? Follow Google instructions

<https://twitter.com/googledocs/status/730087240156643328>



and many more...

# Architectural Changes


Re-done from milestone 1 for efficiency

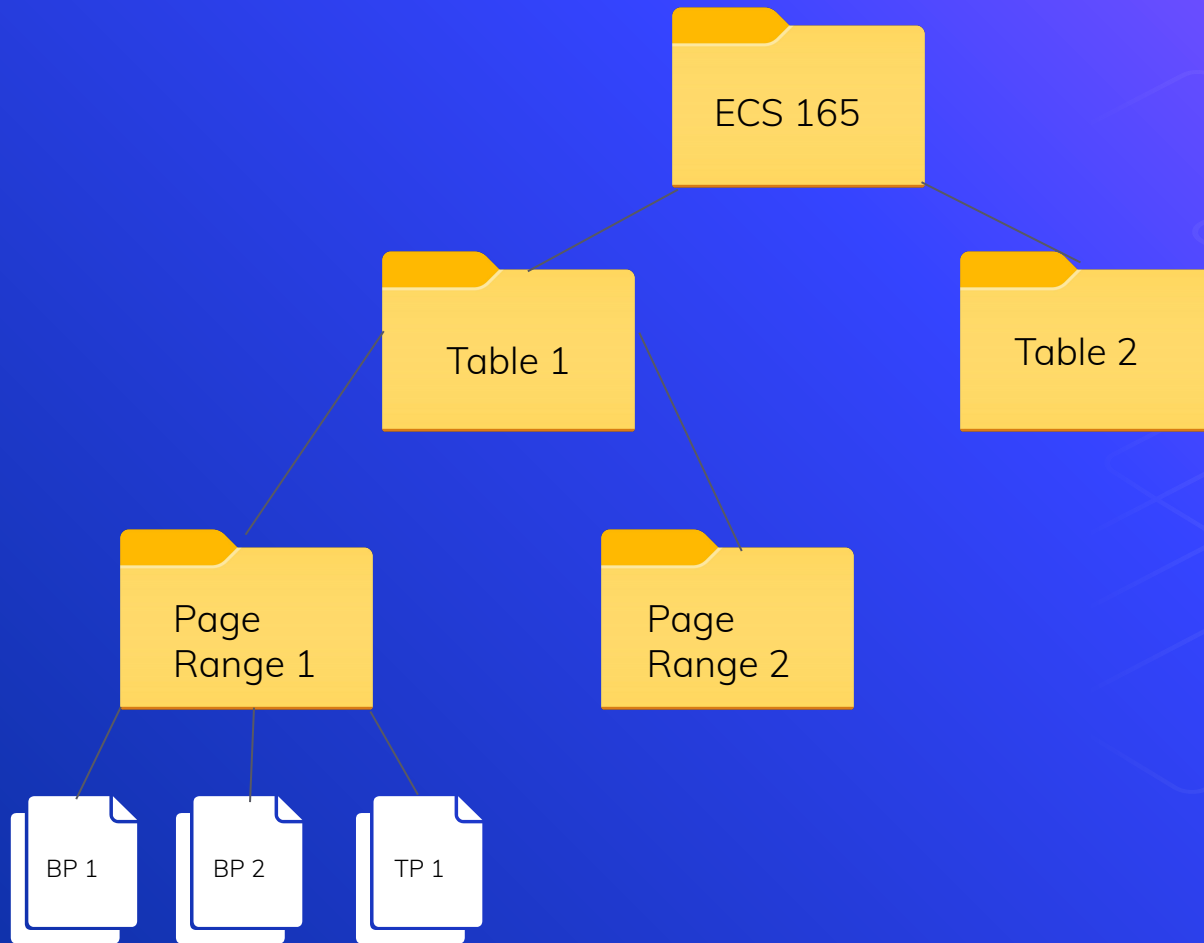


# Writing To Disk



- Write CP object to a file using pickle module





# Data Management - Bufferpool

- ⬡ MetaData
- ⬡ Conceptual Page granularity
- ⬡ Manages interactions between Query and Disk





# Bufferpool - Metadata

- ⬡ Maintains key\_dict - > Key : Path
- ⬡ Stores any other dict based on what is indexed
- ⬡ Keeps track of # PR, BP, RID etc.



# Bufferpool- Evict

- ⬡ LRU Eviction Policy
- ⬡ Only evicts unpinned pages
- ⬡ Cleans up by locally stored keys

```
### Assuming pages stack will be LRU at top of stack (index 0)
def evict(self): #evict a physical page from bufferpool (LRU)
    ### check if value being evicted is pinned
    # Write to disk whatever is at the top of stack
```


```
    i = 0
    temp_cpage = self.conceptual_pages[i]
    while temp_cpage.isPinned:
        i += 1
        if i == len(self.conceptual_pages):
            i = 0
        temp_cpage = self.conceptual_pages[i]
    self.conceptual_pages.pop(i)
```

```
    self.remove_keys(temp_cpage)
    with open(temp_cpage.path, 'wb') as db_file:
        pickle.dump(temp_cpage, db_file)
    # TypeError: file must have a 'write' attribute
```

```
def remove_keys(self, conceptual_page):
    del self.buffer_keys[conceptual_page.path]
```

# Reformatting Query

Including `bufferpool`



# New Delete

- ⬡ If not updated before “delete” by adding a tail page with none in columns
- ⬡ Otherwise: base\_schema->-0
- Update page with empty Columns
- ⬡ Update uses bufferpool

```
# If not updated, add tail page with None
if not updated:
    # Update to add tail page with None
    self.update(key, *[None]*n_cols)
else:
    # Change base schema to all 0's, then
    base_schema = np.zeros(n_cols)
    self.update(key, *[None]*n_cols)

return True
```

```
def update(self, key, *columns):
    # """--New--"""
    columns_to_update = self.colsToUpdate(key, *columns)
    my_base_page = self.get_from_disk(key=key)
    tail_page = self.get_tail_page(my_base_page, key, *columns)
    return self.update_tail_page(my_base_page, tail_page, key, *columns)
```

# New Insert

- ⬡ Restructured to use bufferpool
- ⬡ Uses the bufferpool to perform insertions

Insert

True

False


Bufferpool

Disk

```
if new_base_page:
    self.buffer_pool.meta_data.currbp += 1
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage = self.buffer_pool.createConceptualPage(path, columns)
else:
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage, is_in_buffer = self.in_buffer(path)
    if not is_in_buffer:
```


# New Select

- ⬡ If we are selecting based on keys simply use `key_dict` and same logic as before
- ⬡ Otherwise we have to utilize secondary indexes if available



# New Update


- Fetch the BP from key
- Find or create the TP associated with given key
- Then update like usual





# New Update - Snapshot

- Snapshot upon first update to a specific column in BP
- Copy all query columns from BP into TP



Base Page : [ 45, 21, 32 , 1, 9 ]

Update: [ 17, None, None, 2, None]

Snapshot : [ 45, max\_int, max\_int, 1, max\_int]

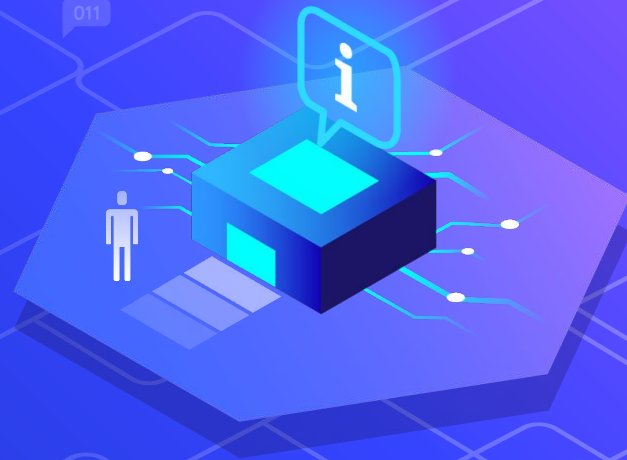


# New Sum




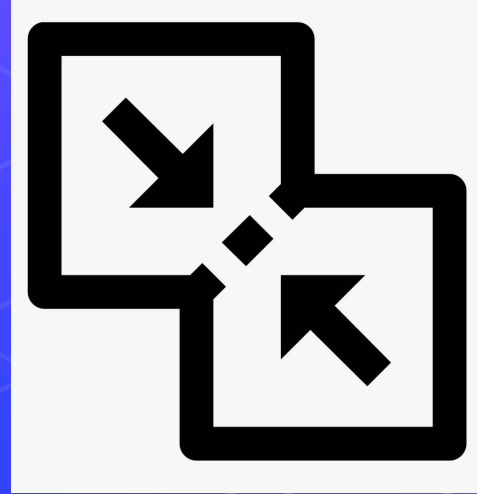
# Data Reorganization

Lazy, Contention Free Merge




# Merge: Criteria

- BP must be read-only 
  - Keep track of qualified BPs in bufferpool
  - Qualified BP is if BP is full
- Merge after every 500 updates, check qualified BP array



# Merge: Process

- Open new thread using Python's Threading library and run in background
- Pull base page and associated tail pages into memory
  - Create a copy of the base pages
- Consolidate updates onto copied base pages
- After merge we change the key\_dict to point to the new consolidated base pages



# Indexing

Create Index & Drop Index




# create\_index()

- Iterate through base pages in disk
- Iterate through records in base page
  - Value: (base\_path, record num)
- Add that value to a index dictionary:
  - (e.g value:[(path\_base\_i, record\_num\_j)])


# drop\_index()

- ⬡ Deletes the dictionary for a given index



# Future Goals & Plans

- Implement Multiple Threads for concurrent use of the L-Store





# New Insert

- ⬡ Restructured to use bufferpool
- ⬡ Uses the bufferpool to perform insertions

Insert

True

False


Bufferpool

Disk

```
if new_base_page:
    self.buffer_pool.meta_data.currbp += 1
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage = self.buffer_pool.createConceptualPage(path, columns)
else:
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage, is_in_buffer = self.in_buffer(path)
    if not is_in_buffer:
```


# New Select

- ⬡ If we are selecting based on keys simply use `key_dict` and same logic as before
- ⬡ Otherwise we have to utilize secondary indexes if available




# New Update

- Fetch the BP from key
- Find or create the TP associated with given key
- Then update like usual



# New Update - Snapshot

- Snapshot upon first update to a specific column in BP
- Copy all query columns from BP into TP

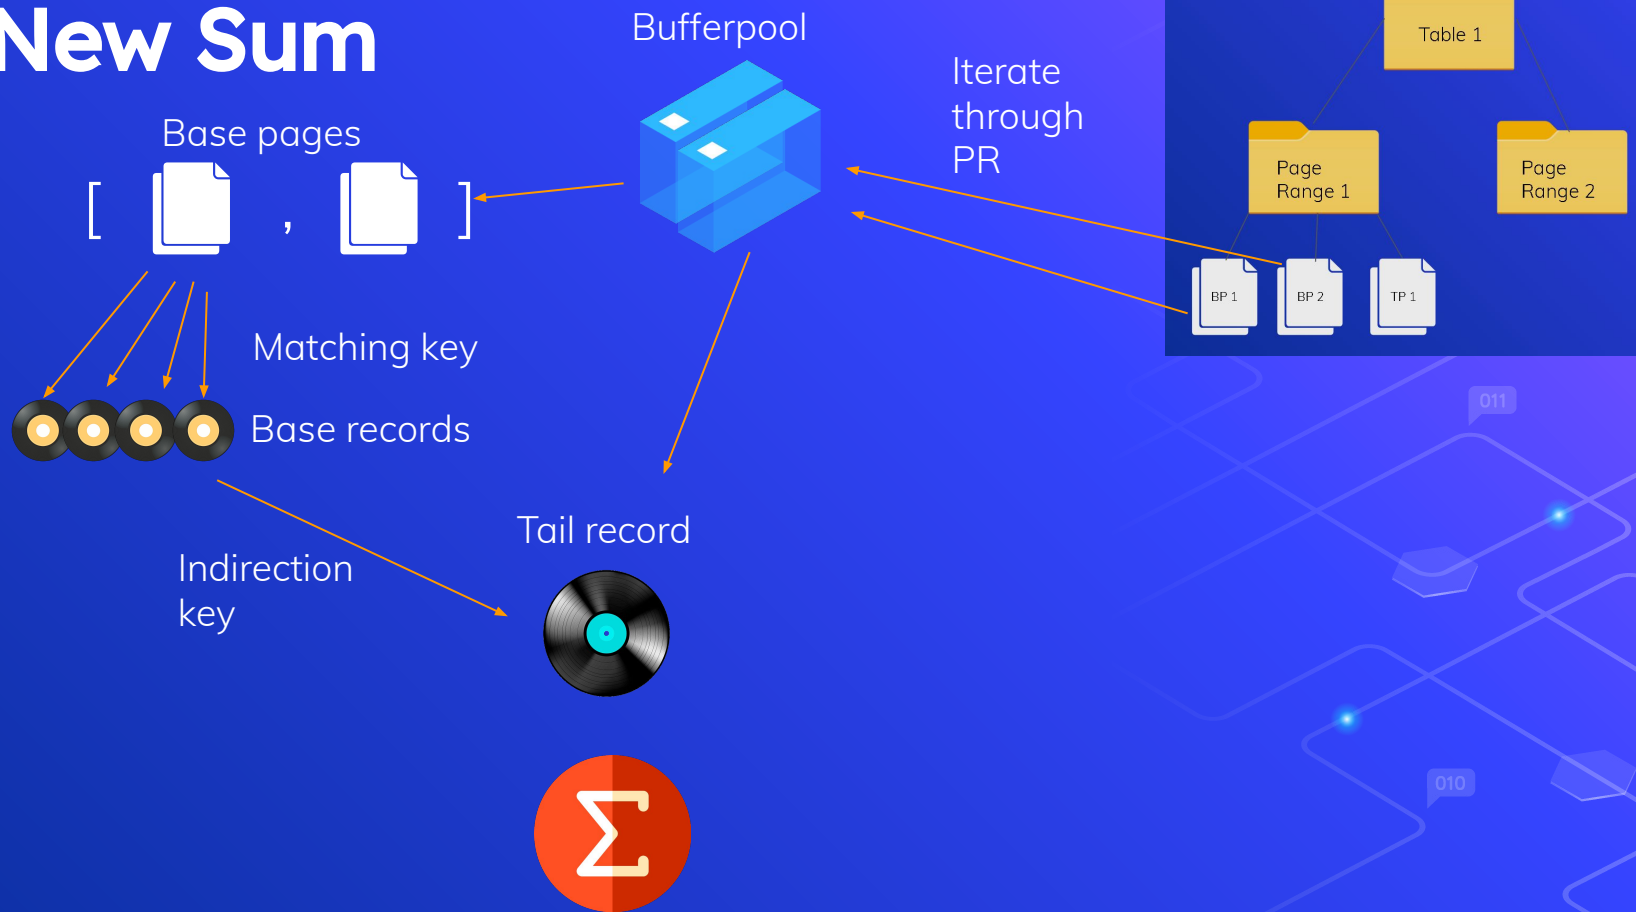


Base Page : [ 45, 21, 32 , 1, 9 ]

Update: [ 17, None, None, 2, None]

Snapshot : [ 45, max\_int, max\_int, 1, max\_int]

# New Sum





# Data Reorganization

Lazy, Contention Free Merge




# Merge: Criteria

- BP must be read-only 
  - Keep track of qualified BPs in bufferpool
  - Qualified BP is if BP is full
- Merge after every 500 updates, check qualified BP array



# Merge: Process

- Open new thread using Python's Threading library and run in background
- Pull base page and associated tail pages into memory
  - Create a copy of the base pages
- Consolidate updates onto copied base pages
- After merge we change the key\_dict to point to the new consolidated base pages





# Indexing

Create Index & Drop Index




# create\_index()

- Iterate through base pages in disk
- Iterate through records in base page
  - Value: (base\_path, record num)
- Add that value to a index dictionary:
  - (e.g value:[(path\_base\_i, record\_num\_j)])


# drop\_index()

- ⬡ Deletes the dictionary for a given index



# Future Goals & Plans

- Implement Multiple Threads for concurrent use of the L-Store



# New Insert

- ⬡ Restructured to use bufferpool
- ⬡ Uses the bufferpool to perform insertions

Insert

True

False


Bufferpool

Disk

```
if new_base_page:
    self.buffer_pool.meta_data.currbp += 1
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage = self.buffer_pool.createConceptualPage(path, columns)
else:
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage, is_in_buffer = self.in_buffer(path)
    if not is_in_buffer:
```


# New Select

- ⬡ If we are selecting based on keys simply use `key_dict` and same logic as before
- ⬡ Otherwise we have to utilize secondary indexes if available




# New Update

- Fetch the BP from key
- Find or create the TP associated with given key
- Then update like usual



# New Update - Snapshot

- Snapshot upon first update to a specific column in BP
- Copy all query columns from BP into TP



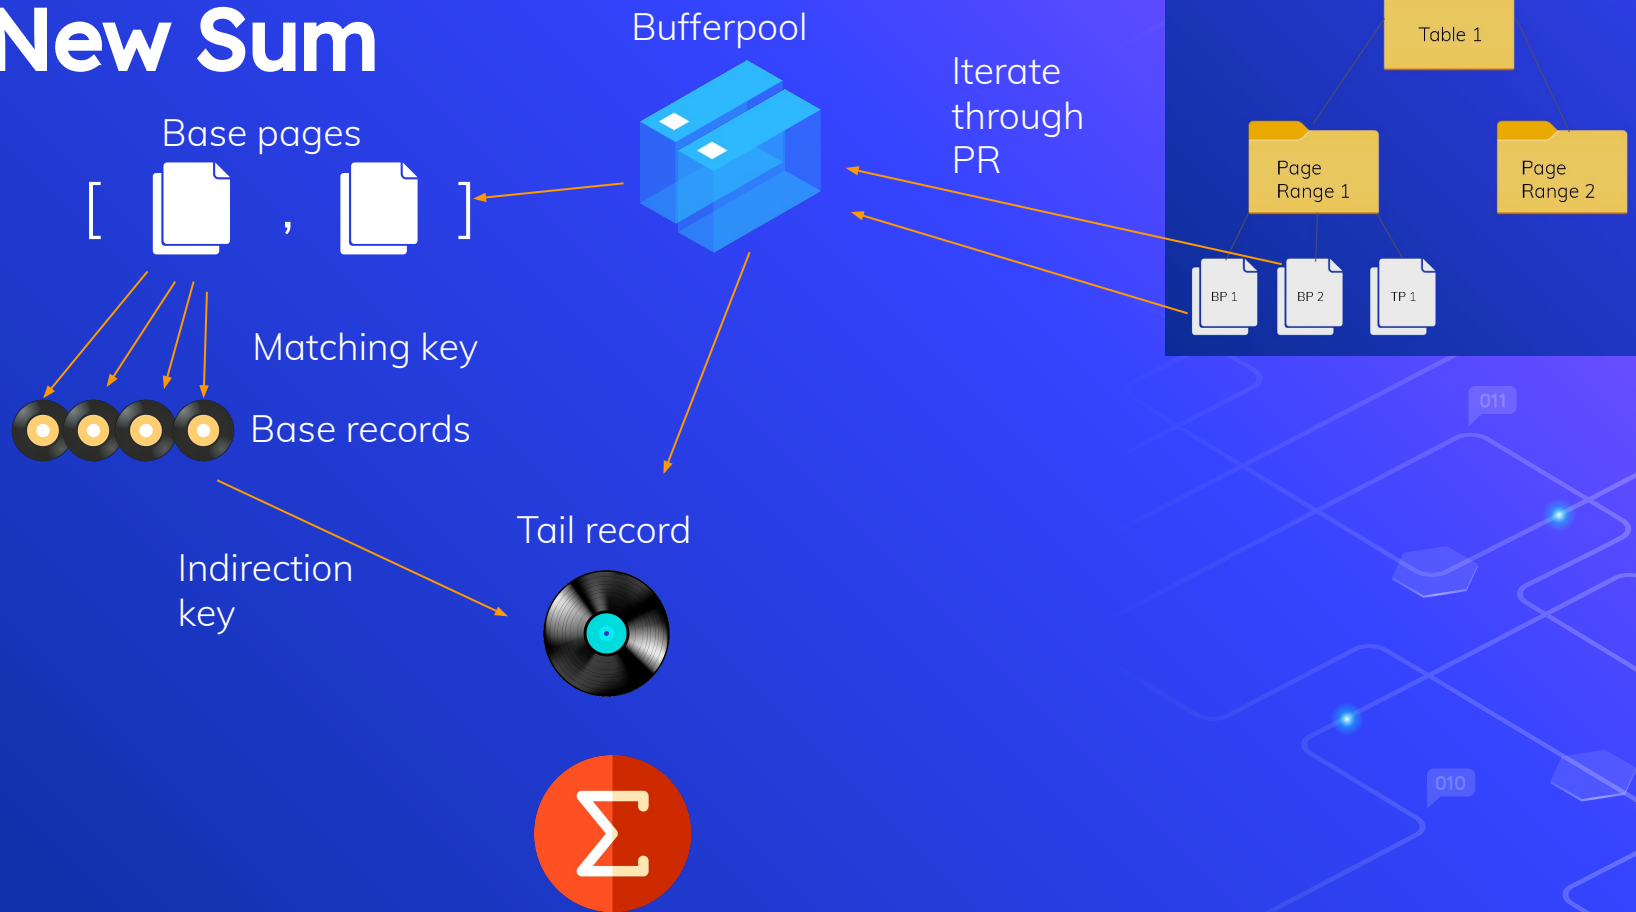
Base Page : [ 45, 21, 32 , 1, 9 ]

Update: [ 17, None, None, 2, None]

Snapshot : [ 45, max\_int, max\_int, 1, max\_int]



# New Sum




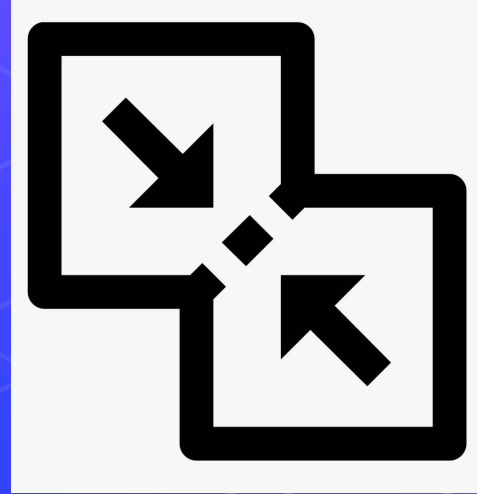
# Data Reorganization

Lazy, Contention Free Merge




# Merge: Criteria

- BP must be read-only 
  - Keep track of qualified BPs in bufferpool
  - Qualified BP is if BP is full
- Merge after every 500 updates, check qualified BP array



# Merge: Process

- Open new thread using Python's Threading library and run in background
- Pull base page and associated tail pages into memory
  - Create a copy of the base pages
- Consolidate updates onto copied base pages
- After merge we change the key\_dict to point to the new consolidated base pages



# Indexing

Create Index & Drop Index




# create\_index()

- Iterate through base pages in disk
- Iterate through records in base page
  - Value: (base\_path, record num)
- Add that value to a index dictionary:
  - (e.g value:[(path\_base\_i, record\_num\_j)])


# drop\_index()

- Deletes the dictionary for a given index



# Future Goals & Plans

- Implement Multiple Threads for concurrent use of the L-Store





# New Insert

- ⬡ Restructured to use bufferpool
- ⬡ Uses the bufferpool to perform insertions

Insert

True

False


Bufferpool

Disk

```
if new_base_page:
    self.buffer_pool.meta_data.currbp += 1
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage = self.buffer_pool.createConceptualPage(path, columns)
else:
    path = './ECS165/' + self.table.name + '/PR' + str(self.buffer_pool.meta_data.currpr) + '/BP' + str(self.buffer_pool.meta_data.currbp)
    cpage, is_in_buffer = self.in_buffer(path)
    if not is_in_buffer:
```


# New Select

- ⬡ If we are selecting based on keys simply use `key_dict` and same logic as before
- ⬡ Otherwise we have to utilize secondary indexes if available




# New Update

- Fetch the BP from key
- Find or create the TP associated with given key
- Then update like usual



# New Update - Snapshot

- Snapshot upon first update to a specific column in BP
- Copy all query columns from BP into TP

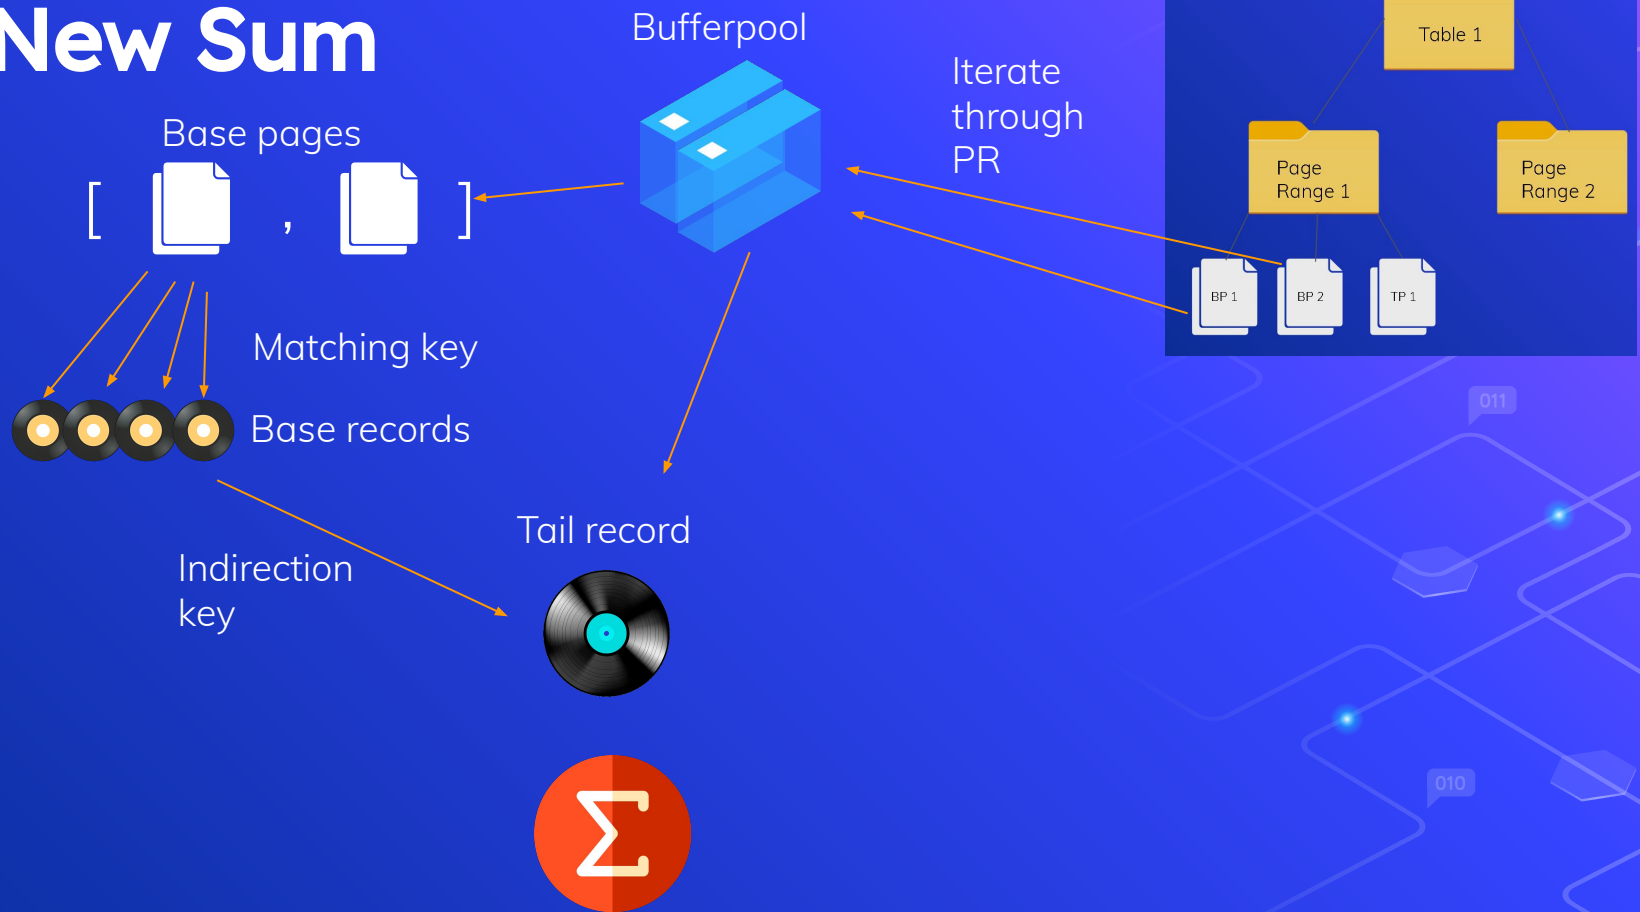


Base Page : [ 45, 21, 32 , 1, 9 ]

Update: [ 17, None, None, 2, None]

Snapshot : [ 45, max\_int, max\_int, 1, max\_int]

# New Sum





# Data Reorganization

Lazy, Contention Free Merge




# Merge: Criteria

- BP must be read-only 
  - Keep track of qualified BPs in bufferpool
  - Qualified BP is if BP is full
- Merge after every 500 updates, check qualified BP array



# Merge: Process

- Open new thread using Python's Threading library and run in background
- Pull base page and associated tail pages into memory
  - Create a copy of the base pages
- Consolidate updates onto copied base pages
- After merge we change the key\_dict to point to the new consolidated base pages





# Indexing

Create Index & Drop Index




# create\_index()

- Iterate through base pages in disk
- Iterate through records in base page
  - Value: (base\_path, record num)
- Add that value to a index dictionary:
  - (e.g value:[(path\_base\_i, record\_num\_j)])

# drop\_index()

- ⬡ Deletes the dictionary for a given index



# Future Goals & Plans

- Implement Multiple Threads for concurrent use of the L-Store

